

# Human-Following Ackermann v2 — Refactor to OAK-D Pro W

**Author:** Senithu Dampegama — Robotics & AI Engineer

**Date:** 23 September 2025

**Version:** v2 (successor to the Arducam/PiNSight-based v1)

## Contents

Human-Following Ackermann v2 — Refactor to OAK-D Pro W.....	1
0) Executive Summary .....	3
1) Background & Motivation .....	3
2) System Architecture (L5→L1) .....	4
3) Hardware Overview .....	4
4) Software Architecture .....	5
4.1 Vision (OAK-D On-Device).....	5
4.2 Tracking (Lightweight DeepSORT) .....	5
4.3 Decision (L5) — FSM.....	5
4.4 Interface (L4) — ASCII Protocol (unchanged) .....	5
5) Repository Layout (v2).....	6
6) Interfaces & Protocols .....	6
6.1 Command Protocol (Jetson → Teensy) .....	6
6.2 Telemetry Protocol (Teensy → Jetson) .....	6
7) Control Parameters (defaults) .....	7
7.1 Steering & Turn Shaping.....	7
7.2 Distance & Motion .....	7
7.3 Robustness.....	7
8) Setup & Runbook.....	8
8.1 Environment .....	8
8.2 STUB (safe, no motion) .....	8
8.3 LIVE (hardware-in-loop).....	8
9) Validation & Observations.....	8
10) Risk & Limitations .....	9

11) Tuning Guide .....	9
12) Future Work.....	9
Appendix A — CLI Flags (selected) .....	9
Appendix B — Serial Protocol Examples .....	10
Appendix C — Safety Checklist .....	10

## 0) Executive Summary

- **Objective:** Upgrade the prior human-following Ackermann robot (v1) to reliably track **one person among many** while reducing Jetson load.
- **Approach:** Offload perception to **OAK-D Pro W** (Myriad X) running a **Spatial MobileNet-SSD**; add a **lightweight IOU DeepSORT** tracker; retain the proven **L5/L4 layered control** with **Teensy 4.1** actuation.
- **Safety & Robustness:** Stationary **SEARCH/RECOVER** (no motion), **serial auto-reconnect** on USB drops, **bounded runtime** flag for tests, capped PWM, and scalable dual preview (RGB + Depth).
- **Results:** Stable person lock with low Jetson usage; tighter right turns via **asymmetric steering gains** and **non-linear mapping**; unchanged ASCII protocol to Teensy; clean STUB→LIVE runbooks.
- **Deliverables:** A reorganized package hf\_ackermann/ ready for reuse, plus this report documenting architecture, parameters, tests, and future work.

## 1) Background & Motivation

**v1 summary.** The previous build transformed a failed RC toy into a research robot: Jetson Orin Nano for perception/decision (L5), Teensy 4.1 for actuation/safety (L3), and a simple ASCII protocol (L4). It followed a single person reliably but **could not lock onto one human** if multiple people entered the frame.

**v2 goals.**

1. **Resolve multi-person ambiguity** via per-person IDs.
2. **Reduce Jetson load** by moving heavy CV to the vision sensor.
3. Preserve the **safe, layered** control split and the exact L4 ASCII protocol to avoid firmware churn.

**Key change:** Replace Arducam/PiNSight with **OAK-D Pro W**; implement a lightweight, IOU-only DeepSORT to keep stable track\_id while minimizing compute.

## 2) System Architecture (L5→L1)

OAK-D Pro W — RGB + StereoDepth + Spatial NN (on-device)

| detections {bbox, conf, z\_mm}



Tracking (IOU DeepSORT) → selected target {track\_id, offx, area, z\_m, quality}



L5 (Jetson) Decision FSM INIT → SEARCH → ACQUIRE → FOLLOW → FAULT

| emits high-level actions (ARC/DRIVE/STOP/CENTER)



L4 (Jetson) ASCII Serial Client → Teensy 4.1 (safety, PWM bursts, stall guard)



Drivers & Actuators (BTS7960 drive, BTS7980 steering)

**What's new in v2** - On-device **Spatial NN** (x,y,z) from OAK-D; Jetson only does tracking + L5.

- **Dual preview:** RGB with overlays + depth false-color (scalable).
- **Stationary SEARCH/RECOVER** (no peeking) for safety and clarity.
- **Asymmetric steering** (right boosted) and **gamma-shaped turn mapping**.
- **Serial auto-reconnect & bounded runtime** for robust testing.

## 3) Hardware Overview

- **Compute:** NVIDIA Jetson (Orin Nano or better).
- **Camera:** OAK-D Pro W (wide FOV), Myriad X.
- **Controller:** Teensy 4.1 (Cortex-M7, 600 MHz).
- **Drivers:** BTS7960 (drive axles), BTS7980 (steering).
- **Power:** Dual 3S LiPo, separate regulators; common ground; shielded USB to Teensy.

*Note:* Keep OAK-D and high-current returns cleanly separated; add ferrites to USB if needed.

## 4) Software Architecture

### 4.1 Vision (OAK-D On-Device)

- **Pipeline:** ColorCamera 1080p + StereoDepth (aligned to RGB) → **MobileNetSpatialDetectionNetwork @ 300×300**.
- **Outputs per detection:** bbox\_px, confidence, label, spatialCoordinates.{x,y,z} in mm.
- **VisionBus payload to L5:** {track\_id, offx∈[-1..1], area∈[0..1], quality, z\_m (m), ts}.
- **Preview (when enabled):** Composite window **[RGB annotated | Depth false-color]**, scalable via --show-scale.
- **Hotkeys:** l lock nearest, r release, q/ESC quit (optional s STOP, c CENTER if wired).

### 4.2 Tracking (Lightweight DeepSORT)

- IOU-only association with max\_age, iou\_threshold, hit\_streak → **sticky IDs** at minimal compute cost.
- Optional heavy wrapper (deep-sort-realtime) available but off by default on Jetson.

### 4.3 Decision (L5) — FSM

States: INIT → SEARCH → ACQUIRE → FOLLOW → FAULT - **SEARCH / RECOVER: stationary**; on entry: STOP + CENTER once; wait for a valid target.

- **ACQUIRE:** quick steering impulses to bring target toward center; short forward nudges if far.

- **FOLLOW:** - **Steering:** non-linear mapping from |offx| with exponent arc\_gamma and a **minimum turn PWM; asymmetric left/right gains & caps** (right boosted).

- **Distance:** prefer **depth z\_m** (target-z), else **area** (target-area ± tol).

- **FAULT:** triggered by L4 disconnects; L5 stops issuing motion, then attempts **auto-reconnect** with backoff; on success: STOP + CENTER and resume SEARCH.

### 4.4 Interface (L4) — ASCII Protocol (unchanged)

- **Commands:** STOP, CENTER, DRIVE F|B <pwm> <ms>, ARC L|R <pwm> <ms> <U|H>.
- **Replies:** OK, DONE <ARC|DRIVE>, ERR <...>, TELEM drive=<...> steer=<...> L=<...> R=<...> t=<...>.
- **Timing:** 115200 baud; read timeout ≈ 0.05 s; cmd timeout ≈ 2 s; DONE timeout ≈ 3 s (or by duration).

- **Robustness:** auto-detect Teensy; **auto-reconnect** on SerialException; resume in SEARCH.

## 5) Repository Layout (v2)

```

hf_ackermann/
  app.py                # CLI entry: STUB/LIVE, vision source,
tuning flags
  control/
    rc_layer4.py        # L4 serial client (ASCII, OK/DONE/TELEM,
reconnect)
    rc_layer5.py        # L5 FSM (stationary SEARCH), steering
shaping, distance logic
  vision/
    vision_oakd_onboard.py  # OAK-D spatial NN + depth + composite
preview
    vision_pinsight_legacy.py  # Legacy fallback (non-spatial)
tracking/
    deepsort_iou.py        # Lightweight IOU tracker (default)
    deepsort_heavy.py      # Optional wrapper for deep-sort-realtime
tests/
    test_l4_smoke.py       # OK/DONE/TELEM smoke
    test_l4_fault.py      # L4 disconnect → auto-reconnect (mocked)
scripts/
    run_stub_oakd.sh       # safe preview, no motion
    run_live_oakd.sh       # hardware-in-loop (wheels-up first!)
requirements.txt, README.md, LICENSE

```

## 6) Interfaces & Protocols

### 6.1 Command Protocol (Jetson → Teensy)

```

STOP\n
CENTER\n
DRIVE F 190 240\n
DRIVE B 140 150\n
ARC L 220 180 U\n
ARC R 220 180 U\n

```

### 6.2 Telemetry Protocol (Teensy → Jetson)

```

OK\n
DONE DRIVE\n
TELEM drive=1 steer=2 L=497 R=1022 t=25090824\n

```

drive/steer are state codes; L/R are sensor counts; t is uptime ms.

## 7) Control Parameters (defaults)

### 7.1 Steering & Turn Shaping

Parameter	Default	Purpose
deadband_x	0.16	Ignore small offsets to avoid chatter
arc_gamma	1.6	Non-linear turn shaping for large errors
pwm_turn_min	140	Enforce decisive minimum turn
k_arc_pwm_left	700	Left turn gain
k_arc_pwm_right	820	<b>Right turn boosted</b>
arc_pwm_cap_left	250	Max left arc PWM
arc_pwm_cap_right	300	<b>Max right arc PWM</b>
arc_impulse_ms_left	180 ms	Left arc duration
arc_impulse_ms_right	200 ms	<b>Right arc duration</b>
steer_bias	0.00	Optional static offset (normalize drift)

### 7.2 Distance & Motion

Parameter	Default	Purpose
target_z_m	(off)	Preferred distance (m) if depth available
target_area	0.18	Fallback distance via bbox area
area_tol	0.04	Area band around target
k_drive_pwm	2200	Forward PWM gain (capped)
max_drive_pwm	190	Safety cap for drive
drive_impulse_ms	240 ms	Forward impulse duration
back_pwm	140	Reverse PWM for back-off
back_impulse_ms	150 ms	Reverse impulse duration

### 7.3 Robustness

Parameter	Default	Notes
search_mode	idle	SEARCH/RECOVER are stationary
--show-scale	0.5–0.6	Smaller preview window
--max-seconds	(off)	Bounded runtime for tests

All are exposed via CLI flags in `app.py` for runtime tuning.

## 8) Setup & Runbook

### 8.1 Environment

```
python3 -m venv .venv && source .venv/bin/activate
pip install -r requirements.txt
```

#### OAK-D probe

```
python - <<'PY'
import depthai as dai
with dai.Device() as d:
    print('OAK-D connected:', d.getDeviceName(), d.getMxId())
PY
```

### 8.2 STUB (safe, no motion)

```
./scripts/run_stub_oakd.sh
# or
python -m hf_ackermann.app --mode STUB --vision OAKD --show 1 --max-seconds 8
```

### 8.3 LIVE (hardware-in-loop)

**Safety:** first runs **wheels-up** on a stand; ensure serial permissions (dialout).

```
./scripts/run_live_oakd.sh
# or
python -m hf_ackermann.app --mode LIVE --vision OAKD --show 1 --max-seconds 10
```

If auto-detect fails: `--port /dev/ttyACM0` (or a udev symlink `/dev/teensy`).

## 9) Validation & Observations

- **Stationary SEARCH:** verified — no ARC/DRIVE while waiting; transitions to ACQUIRE on fresh target.
- **FOLLOW:** decisive turns, especially to the right (boosted gains/caps); depth-preferred distance with smooth back-off when close.
- **Serial robustness:** induced disconnects yield [FAULT] → **auto-reconnect** → STOP + CENTER → back to SEARCH without crash.
- **Multi-person:** lock persists via `track_id`; l/r hotkeys control target selection.

#### Representative telemetry

```
TELEM drive=1 steer=2 L=497 R=1022 t=25090824
[FOLLOW] steer arcright pwm=220
```



TELEM drive=1 steer=2 L=526 R=1021 t=25091024  
[FOLLOW] stale → center once  
...

## 10) Risk & Limitations

- No wheel encoders → open-loop velocity; rely on depth/area and impulses.
- Donor wiring caps drive PWM to 190 for thermal safety; revisit after loom upgrade.
- IOU-only tracker can ID-swap under heavy, long occlusions (lock + stationary SEARCH mitigates most cases).
- Depth noise possible in strong sun/low-texture scenes; consider median filtering or depth caps.

## 11) Tuning Guide

- **Right turns weak?** Increase `--k-arc-right`, `--cap-right`, `--impulse-right`; optionally raise `--pwm-turn-min` and `--arc-gamma`.
- **Over-steer/oscillation?** Increase `--deadband-x` by +0.02 or lower `--arc-gamma` (1.4–1.6).
- **Far off-center?** Increase `--arc-gamma` (e.g., 1.8) and/or `--pwm-turn-min`.
- **Depth-first following:** set `--target-z` 1.2 (example), otherwise use `--target-area/--area-tol`.

## 12) Future Work

- Appearance embeddings (heavy DeepSORT) when compute budget allows.
- Add encoders + basic odometry; consider SLAM for navigation.
- Wiring/driver upgrades to lift PWM caps safely.
- systemd auto-start (STUB by default), simple web telemetry, and a udev rule for `/dev/teensy`.

## Appendix A — CLI Flags (selected)

```
--mode {STUB,LIVE}  
--vision {OAKD,FAKE,PINSIGHT}  
--port /dev/ttyACM0  
--show {0,1}          --show-scale 0.5
```

```
--max-seconds 10

# Steering & distance
--deadband-x 0.16
--arc-gamma 1.6
--pwm-turn-min 140
--k-arc-left 700 --k-arc-right 820
--cap-left 250 --cap-right 300
--impulse-left 180 --impulse-right 200
--steer-bias 0.00
--target-z 1.2
--target-area 0.18 --area-tol 0.04
```

## Appendix B — Serial Protocol Examples

```
> ARC R 220 180 U
< OK
< DONE ARC

> DRIVE F 190 240
< OK
< DONE DRIVE

< TELEM drive=1 steer=2 L=615 R=1022 t=25091724
```

## Appendix C — Safety Checklist

- Start in **STUB**; only then run **LIVE** (wheels-up).
- Use --max-seconds during tests to avoid endless runs.
- Confirm TELEM rate (~5 Hz), caps (max\_drive\_pwm, arc\_pwm\_cap\_\*).
- Verify **auto-reconnect** path by briefly replugging USB (observe FAULT→recover).
- Keep grounds clean; use short, shielded USB.